

Intelligence Operations for Entitlements

The Skunkworks playbook for land-use lobbying firms.

A watered-down — or beefed-up — Cambridge Analytica for local public affairs.

PREPARED FOR Lattice Public Affairs (white-label)
CLASSIFICATION Internal · methodology + legal positioning
COMPANION TO The War Room · live intel surface

S 1 · TL;DR

The pitch in one paragraph.

Ship one signature feature inside the War Room: the **Hearing Eve Brief** — a one-click, day-before-the-hearing intelligence dossier that names every commissioner's likely vote with a confidence score, every speaker likely to testify against the project, every argument they'll raise, and a pre-written rebuttal for each. It is the feature a senior partner cannot unsee, because it converts the firm's most stressful, manual, all-nighter ritual into a single button — and it visibly compounds with every project the firm runs.

Behind it, stack a four-layer intelligence-operations platform that becomes uncopiable through network effects, the way Bloomberg Terminal, Palantir Foundry, and the partisan voter-file co-ops did: (1) cross-project / cross-firm benchmarking via a de-identified data pool; (2) name-level predictive opposition modeling; (3) real-time adversarial intelligence on opposition group formation and funding; (4) elected-official influence mapping and vote tracking.

Most of the highest-value playbook is clearly legal — public records, voter files, donor disclosure, lobbying filings, court records, public social posts, ownership graphs. The aggressive edge — login-gated scraping, cultivated insiders, PI/HUMINT, covert recording — is where the legal and reputational risk concentrates. Ship the legal core. Keep the gray-area work as partner-referral, never native.

§ 2 · KEY FINDINGS

Five things that matter.

1. There is a proven public comparable for the adversarial-intelligence layer.

10a Labs' **Data Center Watch** publicly catalogs activist groups organizing against data center construction — 188+ organized opposition groups across 40 states, with reporting tying \$156 billion in blocked or delayed projects to local opposition, moratoria, and litigation. It is built on LLMs + human analysts + OSINT. The market wants this kind of product, and the recipe works.

2. Opposition is often not what it appears — and exposing that is a killer feature.

Holland & Knight's 2015 study of every CEQA suit filed 2010–2012 found recognized environmental advocacy groups comprised only 13% of CEQA petitioners. Sixty-four percent were individuals or local "associations" with no prior environmental record. Competitor- and union-funded greenmail, RICO cases against rival hotel developers, and bounty-hunter law firms are well-documented. A feature that traces opposition funding and unmasks astroturf sits on solid public-records ground.

3. Name-level prediction is technically real but accuracy-bounded.

Commercial voter files (L2 ~217M records; TargetSmart VoterBase ~208M voters + \$6.7B in contribution history) plus consumer enrichment (i360, Acxiom, Experian) make individual propensity scoring real. But published research caps accuracy modestly — vote-choice models in the high 60s%; turnout models with Matthews correlations around 0.39. Sell a prioritized *likely-objector list*, not a deterministic prophecy.

4. The compounding moat is real and has precise precedents.

Palantir Foundry explicitly supports multi-org ecosystems that pool participant data for modeling, governed by data-sharing agreements and anonymization. The partisan co-ops (Democratic Data Exchange, Republican Data Trust / i360) are blind exchanges with give-to-get rules so non-coordinating entities legally pool data. Bloomberg shows where this ends: a single seat at ~\$32K/year in 2026, up ~60% since 2010, because "switching costs are high."

5. The legal lines are now well-mapped.

Scraping *public* web data does not violate the CFAA (hiQ v. LinkedIn, 9th Cir. 2022). But using fake accounts or scraping behind logins is breach of contract — hiQ ultimately took a \$500K judgment plus permanent injunction. Voter-file commercial use is prohibited in 37 states; New Mexico's HB4 (2023) makes misuse a fourth-degree felony. Covert recording is a felony in ~11–13 two-party-consent states (CA, FL, MA, PA, etc.). These are bright lines you can engineer around.

S 3 · THE FOUR INTELLIGENCE LAYERS

What stacks into the moat.

A. Cross-project / cross-firm intelligence (the network-effects layer)

Every entitlement battle produces structured exhaust — opposition profile, objection timeline, commissioner behavior, revision count, time-to-decision, response effectiveness, outcome. Pool and de-identify it across the firm's projects (and ideally across firms) and you ship:

- **Benchmarking:** "Projects with this opposition profile passed at X%. Your risk score is Y."
- **Pattern detection:** "When objection volume spikes on day 14, Planning Commission rejection probability rises unless addressed."
- **Tactic library:** "When this opposition archetype appeared, the winning response was Z."
- **Comparable timeline / pricing:** "Similar projects took 9–14 months and N revisions."

How to make it legally defensible

Contract construct. A de-identified / aggregated-data clause in the IP section of the MSA, modeled on standard SaaS practice. Cross-firm pooling borrows the political co-op structure — a neutral blind-exchange intermediary plus a give-to-get rule, exactly how DDx and Data Trust / i360 let competitors share. **Anonymization standard.** Clean-room techniques (TEEs, k-anonymity thresholds, differential privacy) with a minimum-aggregation floor before any benchmark is shown.

Position

GREEN · Provided anonymization meets CCPA/GDPR de-identification and the MSA clause is clean. Network effect is slow but uncopyable.

B. Name-level predictive opposition modeling

Data sources combine: voter files (L2, TargetSmart, Catalist, i360, NGP VAN) for registration / vote history / modeled partisanship; donor data (FEC, state campaign finance, OpenSecrets) for ideology and issue intensity; municipal public-comment archives via FOIA bulk extraction — *the single strongest predictor*; property records (Regrid) for proximity and ownership; HOA / neighborhood-association membership; public LinkedIn / professional data. Enrichment via Acxiom, Experian, Resonate, Datavant.

Architecture: logistic regression or gradient-boosted classifier on prior-participation indicators (did they comment / appeal before?), proximity, tenure, partisanship, plus social-graph propagation from known organizers. **Realistic output:** a ranked list with calibrated probabilities, not certainty.

Legal lines — must be enforced in code

- Voter-file commercial use is barred in 37 states. Entitlement lobbying may or may not qualify as a permitted "political" purpose. Build a state-by-state rules engine that gates which fields can be used where.
- CCPA / CPRA and other state privacy laws create access, deletion, and opt-out obligations on individual profiles.

- Scraping Nextdoor / Facebook / HOA directories behind logins is ToS breach plus potential CFAA exposure via fake accounts. Keep ingestion to genuinely public data.

Position

AMBER · Legal in most states with the right purpose and public-only data. Turns **RED** the moment it uses login-gated scraped data or runs in a commercial-use-prohibited state.

C. Adversarial intelligence — who is organizing against the project

Capabilities: real-time detection of opposition group formation (Nextdoor, Facebook groups, X, local press, Substack); identifying ringleaders and their networks; mapping the emerging argument set; a cross-jurisdiction "serial objector" database; and funding traces (is opposition bankrolled by a competing developer, a union, an industry rival?).

Tools / stack: social listening (Brandwatch, Meltwater, Talkwalker, Sprinklr — \$15K–\$150K+/yr); OSINT (Babel Street, ShadowDragon's SocialNet across 200+ sources, Maltego, Recorded Future); ownership / funding networks (Sayari Graph — 455M+ companies, source-linked beneficial-ownership). The Data Center Watch model proves LLM + analyst + OSINT works.

Position

GREEN for public posts, public records, court filings, lobbying / campaign-finance disclosures, ownership graphs. **AMBER** / **RED** for login-gated scraping, infiltrating private groups with fake accounts, any covert recording.

D. Real-time elected-official pulse and influence mapping

Capabilities: commissioner / council-member profiles (voting history, statements, donor base, professional / family / business ties, recent appearances, public social activity); influence graphs (whose calls do they take — chief of staff, neighborhood leader, union, business partner); real-time pulse from social / press / donor filings / agenda patterns to detect shifts before public statements; per-member "what works" coaching generated from cross-project learnings.

Comparable products: Quorum, FiscalNote / CQ, Phone2Action / Capitol Canary, Bloomberg Government, POLITICO Pro, Axios Pro, LegiStorm. **The gap:** none are land-use / local-entitlement specific. None fuse the influence graph with project-level outcome data. That white space is Lattice's wedge.

Position

GREEN for public records, disclosures, public statements, first-party contact logs. **AMBER** for inferred personal / family / business relationships (accuracy + defamation risk — require sourcing). **RED** for surveilling officials, covert recording, cultivating insider sources in offices.

§ 4 · THE SIGNATURE FEATURE

Hearing Eve Brief.

One click, the night before a hearing, produces:

- **Per-commissioner vote call** with a confidence score and the 2–3 levers most likely to move that specific person (from cross-project learnings).
- **Likely-speaker roster** — who will show up to testify against, with each person's probable argument.
- **Pre-written rebuttal** for every predicted argument, grounded in the RAG layer over the Comp Plan and zoning code.
- **"Last time you faced this HOA / commissioner, here's what worked"** — cross-project memory card.

Why it wins the room: it targets the partner's single most painful, manual ritual — the all-nighter before a hearing — and it gets visibly smarter every project. The compounding moat becomes tangible at first sight.

Runner-up wow features for the roadmap: Time-to-Approval Calculator ("11 months vs. 18-month regional average — ~\$3M saved in carry"); Opposition Early-Warning ("new group forming, escalating, here's who"); Litigation-Risk Score (which approvals get challenged in court).

§ 5 · RISK MATRIX

Feature → exposure → framing.

FEATURE	LEGAL	REPUTATIONAL	FRAMING / MITIGATION
Cross-project de-identified benchmarking	GREEN	LOW	"Industry benchmarking & best practices." MSA clause + k-min + opt-out.
Cross-firm pooled intelligence	AMBER	MED	Neutral blind-exchange + give-to-get; antitrust counsel review.
Name-level objector prediction (public + licensed data)	AMBER	HIGH	"Stakeholder mapping & outreach prioritization." State rules engine; public data only.
Adversarial intel on public groups / funding	GREEN	MED	"Due diligence & opposition transparency." Public records + ownership graphs.
Serial-objector cross-jurisdiction database	AMBER	HIGH	"Repeat-participant registry"; sourcing + accuracy controls; no derogatory labels.
Elected-official influence graph	AMBER	HIGH	"Compliance & relationship mapping." Source every edge; public data.
Login-gated scraping (Nextdoor / FB / HOA)	RED	HIGH	Do not ship as native. Public data only.
Covert recording / pretext interviews	RED	SEVERE	Never native. Referral to licensed counsel / PI who owns compliance.
Cultivated insiders / defectors / HUMINT	RED	SEVERE	Never productize. Partner-referral only; client-directed.

S 6 · BUILD DIFFICULTY & PRICING

What ships when. What it costs.

Build-difficulty ranking

Easy to demo now (weeks)

Hearing Eve Brief MVP (LLM over existing RAG + sentiment + heatmap data); social-listening-fed opposition early warning; commissioner public-profile cards; time-to-approval calculator (heuristic first).

Medium (quarters)

Name-level objector model (gated on data licensing); funding / ownership traces via Sayari; litigation-risk score.

Hard / slow but defensible (the moat)

Cross-project benchmarking with statistical validity (needs data volume); cross-firm pooled exchange (needs legal structure + trust + density); per-commissioner predictive vote model; serial-objector cross-jurisdiction database.

Pricing implications

Category shift: "site builder" (low-thousands / year) → "intelligence operations platform" (mission-critical, high-margin). Market anchors: Bloomberg Terminal ~\$32K / seat / yr; FiscalNote / Quorum ~\$10K–\$30K+ / yr; Babel Street / ShadowDragon six-figure annual; social listening \$15K–\$150K+ / yr.

Target a platform subscription per firm in the low-to-mid six figures / year for the full intel-ops suite, plus a per-project intelligence upcharge of \$50K–\$150K for a contested entitlement's Win Room + adversarial-intel package.

Justify it against **carry cost**: practitioners cite \$50K / month on a value-add reposition and \$500K / month on a large ground-up project during entitlement. Save a few months and the platform pays many times over. The buyer shifts from project lead to **managing partner** — sell the moat ("every project makes the next one easier") and the competitive narrative ("your competitors are flying blind").

§ 7 · THE PITCH

Five slides to a senior partner.

1. **"The night before the hearing."** Open on the partner's pain: the all-nighter, the guesswork, the speaker who blindsides you. One slide, visceral.
2. **The Win Room.** Live demo of the Hearing Eve Brief: per-commissioner vote calls, predicted speakers + arguments, auto-rebuttals, "last time this HOA" memory. *This is the unsee-able moment.*
3. **The intelligence stack underneath.** Four layers — cross-project benchmarking, name-level prediction, adversarial / funding intel, elected-official influence graph — with the legal-color coding shown openly (green core, amber options, red excluded). Signals sophistication and defensibility.
4. **Why it compounds (the moat).** Bloomberg / Palantir / political-co-op analogy: every project the firm runs makes every future project's model sharper. Cross-firm pooling — blind exchange, give-to-get — makes Lattice the category's permanent infrastructure. "Citizens United for entitlements."
5. **The math + the ask.** Carry-cost ROI (months saved × \$50K–\$500K / month) vs. price; pricing tiers; framing: ship the legal core, refer the gray-area work to compliant partners. Close on exclusivity / first-mover in their market.

§ 8 · CAVEATS

What to keep honest about.

Accuracy. Individual-level political-behavior prediction tops out modestly (~66% vote-choice in published models; turnout MCC ~0.39). Sell ranked prioritization, not certainty. Over-promising invites client distrust and a Cambridge-Analytica-style backlash.

State-by-state variance is the central engineering constraint — not an afterthought. Voter-file commercial-use bans (37 states), two-party recording felonies (~11–13 states), and divergent FOIA bulk-access rules must be encoded as hard gates in the platform.

Reputational tail risk is asymmetric. Definers (Facebook), Cambridge Analytica, and Black Cube show that even legal-adjacent intelligence work, once public, can become an existential brand event. The aggressive HUMINT / scraping capabilities are best left as client-directed referrals to licensed PIs / counsel who carry the compliance burden — never as native Lattice features.

Antitrust / info-sharing optics of cross-firm pooling warrant counsel review before launch, even though the political co-ops establish a workable template.

Several market figures (social-listening and OSINT contract values, Data Center Watch totals) come from vendor or advocacy sources and should be treated as directional.

*Lattice · Skunkworks Dossier · Companion document to the War Room.
Restricted — partner eyes only.*